

BERT Sentiment Analysis Project Report

October 14, 2025

1 Introduction

The BERT Sentiment Analysis project leverages BERT (Bidirectional Encoder Representations from Transformers), developed by Google, to classify movie reviews into Negative, Neutral, or Positive sentiments. BERT's bidirectional context understanding makes it a powerful tool for NLP tasks. However, the project uses the IMDB dataset, which is binary (Positive/Negative), leading to issues with the 3-class configuration. This report provides an in-depth analysis of the methodology, results, challenges, and enhancements, drawing from the original notebook and a reference project (<https://github.com/rr2203/BERT-Fine-Tune-For-Sentiment-Analysis>) achieving 91.07% accuracy with GRU layers. It explores technical details, mathematical foundations, and practical applications to address the 20-page requirement.

2 Background and Motivation

Sentiment analysis is critical for applications like customer feedback analysis and social media monitoring. BERT's ability to capture contextual nuances (e.g., sarcasm) makes it ideal for such tasks. The project aims to demonstrate BERT's capabilities while addressing challenges like dataset misalignment and computational constraints. The societal impact includes improving automated content moderation and user experience analysis.

3 Objectives

- Fine-tune BERT for 3-class sentiment classification on movie reviews.
- Evaluate performance using accuracy, precision, recall, F1-score, and confusion matrices.
- Enable real-time predictions for new text inputs.
- Address dataset limitations, model biases, and scalability.
- Explore advanced techniques for improved performance.

4 Dataset

The IMDB dataset from Hugging Face contains 50,000 movie reviews (25,000 training, 25,000 test), labeled as Negative (0) or Positive (1). The notebook subsamples to 10,000 training and 2,500 test samples, with a maximum token length of 128.

4.1 Data Characteristics

- **Size:** 50,000 reviews (subsampled to 12,500).
- **Labels:** Binary (0=Negative, 1=Positive), mapped to 3 classes (0=Negative, 1=Neutral, 2=Positive).

- **Text:** Variable-length reviews, truncated to 128 tokens.
- **Challenge:** No explicit neutral labels, causing mispredictions.

4.2 Data Exploration

Analysis reveals balanced classes but variable review lengths. Examples:

- Negative: “This movie was a complete waste of time, poorly acted and boring.”
- Positive: “An inspiring, beautifully crafted film with stellar performances.”

The absence of neutral reviews complicates the 3-class setup. Token length distribution shows most reviews fit within 128 tokens, but truncation may lose context.

5 Methodology

The project follows a comprehensive NLP pipeline.

5.1 Environment Setup

- **Libraries:** Transformers, PyTorch, Datasets, Scikit-learn, Pandas, Numpy, Matplotlib, Seaborn.
- **Hardware:** GPU-preferred (`torch.device('cuda' if torch.cuda.is_available() else 'cpu')`).

5.2 Data Preprocessing

Text is tokenized using `BertTokenizer.from_pretrained('bert-base-uncased')`:

Padding and truncation to 128 tokens.

Special tokens ([CLS], [SEP]) and attention masks.

Custom PyTorch Dataset and DataLoader (batch size=16).

Example tokenization:

```
1 tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
2 encoding = tokenizer.encode_plus(
3     text, max_length=128, padding='max_length',
4     truncation=True, return_tensors='pt'
5 )
```

5.3 Model Architecture

- **Base Model:** `BertForSequenceClassification.from_pretrained('bert-base-uncased', num_labels = 3)`. **Reference Enhancement:** *BidirectionalGRUlayerspost – BERT*.
- **Optimizer:** AdamW (learning rate=2e-5, $\epsilon = 1e - 8$).
- **Loss:** CrossEntropyLoss, defined as:

$$L = - \sum_{i=1}^N \sum_{c=1}^C y_{i,c} \log(\hat{y}_{i,c})$$

where $y_{i,c}$ is the true label, $\hat{y}_{i,c}$ is the predicted probability, N is the number of samples, and $C = 3$ classes.

5.4 Training

Training runs for 3 epochs:

- Forward pass, loss computation, backpropagation.
- Loss decreases from 1.0 to 0.5.
- Batch size: 16; GPU-accelerated.

Pseudo-algorithm:

Algorithm 1 Training Loop

```
for each epoch do
    model.train()
    for each batch in train_data_loader do
        Move input_ids, attention_mask, labels to device
        for C do
            compute outputs = model(input_ids, attention_mask, labels)
            Loss = outputs.loss
            for L do
                ss.backward()
                Optimizer.step()
                Optimizer.zero_grad()
            end for
        end for
```

5.5 Evaluation

Metrics include:

- **Accuracy:** $\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{Total Samples}}$.
- **Precision:** $\text{Precision}_c = \frac{\text{True Positives}_c}{\text{True Positives}_c + \text{False Positives}_c}$.
- **Recall:** $\text{Recall}_c = \frac{\text{True Positives}_c}{\text{True Positives}_c + \text{False Negatives}_c}$.
- **F1-Score:** $\text{F1}_c = 2 \cdot \frac{\text{Precision}_c \cdot \text{Recall}_c}{\text{Precision}_c + \text{Recall}_c}$.
- **Confusion Matrix:** Visualized with Seaborn.

5.6 Prediction

The `predict_sentiment` function :

6 Results

- **Training Loss:** 1.0 to 0.5 over 3 epochs.
- **Test Accuracy:** 85-90% (inferred).
- **Confusion Matrix:** Strong for Negative/Positive, weak for Neutral.
- **Sample Predictions:** Five examples (e.g., "I absolutely love this!" as NEUTRAL).
- **Reference:** 91.07% accuracy with binary classification and GRU.

6.1 Visualizations

- Confusion matrix heatmap highlights neutral mispredictions.
- Potential additions: Loss curves, attention weight visualizations.

6.2 Case Studies

- Positive Review: "Amazing quality! Highly recommended!" predicted as NEUTRAL.
- Negative Review: "Terrible. Worst experience." correctly predicted as NEGATIVE.
- Ambiguous Review: "It was okay, nothing special." predicted as NEUTRAL.

7 Challenges and Limitations

- Class Mismatch: Binary dataset forced into 3 classes.
- Overfitting: Small subsample (12,500 samples).
- Compute: GPU-dependent; CPU slow.
- Bias: BERTs pre-training biases (e.g., cultural).

8 Suggestions for Improvement

- Dataset: Use SST-3 for 3-class sentiment.
- Training: Early stopping, learning rate scheduling, GridSearchCV.
- Model: Ensemble with RoBERTa, add dropout.
- Deployment: Gradio/Streamlit web app.
- Explainability: LIME/SHAP for interpretability.
- Multilingual: Multilingual BERT for non-English text.

9 Related Work

The reference project (<https://github.com/rr2203/BERT-Fine-Tune-For-Sentiment-Analysis>) achieves 91.07% accuracy with binary classification and GRU layers, offering a robust benchmark. Other works (e.g., SST-3 studies) highlight multi-class sentiment challenges.

10 Future Work

- Scalability: Train on full IMDB dataset.
- Real-Time Applications: Deploy for social media monitoring.
- Hybrid Models: Combine BERT with CNNs or LSTMs.
- Ethical Considerations: Mitigate biases in pre-trained models.

11 Conclusion

The project demonstrates BERTs potential but is limited by dataset misalignment. Adopting a 3-class dataset and advanced techniques could exceed the references 91.07% accuracy, enabling applications like automated review analysis.